

# CPS: TTP Option: Medium: Multi-objective Control of Catoptric Systems

ROGER CHAMBERLAIN, CHANDLER AHRENS, CHRIS GILL

Dept. of Computer Science and Engineering, School of Engineering and Applied Science

College of Architecture, Sam Fox School of Design & Visual Arts

Washington University in St. Louis

Email: [roger@wustl.edu](mailto:roger@wustl.edu)

## 1. INTRODUCTION

The energy consumption due to buildings (both residential and commercial) is estimated to be 20% to 40% of the total energy usage in developed countries [24], and lighting and heating are two significant components of this energy consumption [18]. Natural light (i.e., sunlight) is a readily available resource that can contribute to both the illumination and heating of structures, yet in the vast majority of circumstances, its use is limited to passive modalities. For example, daylighting (the use of natural light for illumination) design is dominated by passive window positioning and configuration [31] rather than active control mechanisms (see [17] for the few counterexamples).

**FIXME: What are we doing?**

Figure 1a is an image of a prototype catoptric (mirror) surface (called AMP) that was designed, fabricated, and installed through an undergraduate architecture studio taught by C. Ahrens. The installation redirects light from gable ends of an existing building into the darker recesses of the atrium to create better natural lighting where it is desired. In this installation, the mirror positions are fixed.



**Figure 1: Catoptric system prototypes. (a) AMP, TRex building, St. Louis, MO. (b) FIXME: name?, Steinberg Hall, St. Louis, MO.**

In the next generation of this system, which is currently under construction, the mirrors are under active, 2-axis, microprocessor-based control and therefore can be pointed in different directions dynamically as desired over time.

**FIXME: Where (if at all) should we define Open-Source Architecture (OSArc) as distinct from open-source software? How do the notions of OSArc integrate with what we are doing? (This is a Chandler question.) Is [27] a good citation (I just found it on wikipedia)? How about one or more things that Chandler has written?**

**FIXME: Articulate specific research questions below.**

This research will investigate the following questions:

1. *What are the qualitative and quantitative benefits that can be achieved for building daylighting and thermal management through the use of catoptric systems?*

Issues within this question include the ability to articulate the benefits clearly and to quantify them effectively. Clearly, we are in the domain of multi-objective control, so the relationship between the competing goals must also be articulated and quantified. We intend to investigate the use of Markov Decision Processes (MDPs) as an approach to the multi-objective control problem.

2. *How do we provide for the safety, reliability, maintainability, and continued efficacy of these systems?*
3. *Can we design abstractions that encapsulate sub-systems for effective reuse?* Ultimately, we would like to generalize the above into abstractions that can be leveraged more broadly for arbitrary cyber-physical systems development.

**FIXME:** Brief description of who we are and what we've done.

## 2. BACKGROUND AND RELATED WORK

**FIXME:** Describe first two installations.

**FIXME:** Literature review [3, 5, 11, 17, 20].

**FIXME:** Describe MDP-based optimization and our history with it.

Markov Decision Processes (MDPs) [26] represent a general approach to modeling optimization problems and have been applied in a diverse set of application areas [32]. Examples include robotics [1], economics [4], experiment design [16], medical decisions [2], manufacturing [33], agriculture [19], and our own group's use in scheduling [15, 30] and wireless spectrum management [22].

In this proposal we adopt the definition used by Glaubius et al. [15] of a (discrete-time) Markov decision process as a 5-tuple  $(\mathcal{X}, \mathcal{A}, T, R, \gamma)$ , with **states** designated as  $\chi \in \mathcal{X}$ , **actions** designated as  $a \in \mathcal{A}$ , and a transition system,  $T$ , which gives the probability  $P_T(\chi' | \chi, a)$  of transitioning from state  $\chi$  to state  $\chi'$  on action  $a$ . The reward function  $R(\chi, a, \chi') \in \mathbb{R}_{\geq 0}$  describes the reward that accrues when transitioning from state  $\chi$  to state  $\chi'$  via action  $a$ , under a discount factor,  $\gamma$ , to ensure convergence of the long term reward.

## 3. RESEARCH DESCRIPTION

### 3.1. Question 1 – What are the qualitative and quantitative benefits that can be achieved for bulding day-lighting and thermal management through the use of catoptric systems?

There is ample literature that documents the benefits of natural light in human-occupied spaces [20, 21, 25], yet the thermal effects can also be significant [6] (i.e., too much sunlight can increase temperature to an unacceptable degree).

**FIXME:** We will explore the use of MDPs for multi-objective control of catoptric systems.

**FIXME:** Describe our three previous uses of MDPs, and how the value-optimal solution can be estimated heuristically in each case. Chris, can you take a first crack at this?

Our prior research has used Markov Decision Process models [15] to generate resource management policies off-line [12] for non-preemptive sharing of a resource between multiple purposes at once on-line. For example, a meter-tall robot's camera (oriented by a pan-tilt unit similar to the ones we propose to use in our multi-mirror catoptric installations) may be directed downward to identify wire-frame chairs and other obstacles to navigation that other sensors on the robot may have difficulty detecting, or it may be directed upward to identify faces of people at a reception whose images it can then capture. Given distributions of the durations of intervals during which the camera would need to remain pointed in a given direction to complete an individual task, standard policy iteration techniques then can be used to generate run-time policies that in expectation maximize an objective such as adherence to a strictly proportional allocation of the resource over time [15], or even a more general definition of the utility of completing the different tasks at particular times [30]. We also showed that when different distributions of task completion intervals can occur in different modes (e.g., when a robot moves from room to room), it is possible to learn on-line what mode the system is in, or if the mode is known what the distributions are, but not both [13].

However, policy iteration is exponentially expensive, and even the memory requirements to store complete policies for on-line use may be prohibitive in resource-limited systems. We therefore focused next on the policies that were being generated from the models, and discovered consistent structure in those policies that allowed a reasonable heuristic approximation. For simple proportional sharing, a single geometric partition of a simplex could be calibrated to encode the appropriate policy accurately [14]. For utility-based resource sharing multiple disjoint heuristics were needed but the most effective one to use was clearly defined by problem parameters [29].

As a further illustration both of the applicability of MDP-based policy iteration to generate effective resource management policies, we applied similar techniques to manage a much different resource: the transmission spectrum in wireless networks [23]. Although the semantics of that resource differed radically from the pan-tilt camera, the MDP models were reasonably similar. We extended the basic model to include modulation as well as admission decisions, discovered and characterized common structure among the policies that were generated, and again obtained efficient and effective heuristic policies for on-line use [22].

**FIXME:** Describe our approach to applying MDPs to catroptic system control. This includes quantification of benefits, formulation of a combined objective, identifying the control degrees of freedom, encapsulating all of the above in an MDP framework, exploring the MDP state space, seeking to find a heuristic that approximates the value-optimal solution.

### 3.2. Question 2

**FIXME:** The interesting thing here is how properties like safety are dependent on context, e.g., we want concentrated natural light when harvesting energy but not when illuminating a human-occupied space.

### 3.3. Question 3

**FIXME:** We will investigate the viability and utility of two candidate abstractions: direct mirror control and MDP control.

### 3.4. Intellectual Merit

The intellectual contributions of this project are **FIXME: describe summary of intended intellectual merit** [10].

## 4. EVALUATION/EXPERIMENTATION PLAN

**FIXME:** Two sites identified:

1. **BECS Technology, Inc.**, 10818 Midwest Industrial Dr., St. Louis, MO. **FIXME:** BECS Technology is a small manufacturer of electronic control systems in a number of markets (e.g., agriculture, aquatics, refrigeration). We will have access to the HVAC system in their building.
2. **VelociData, Inc.**, 10425 Old Olive Stree Rd., St. Louis, MO. **FIXME:** VelociData is located in the recently announced **39 North** innovation district, which has the Danforth Plant Science Center, Monsanto, Bio Research & Development Growth Park, and Heliz Center Biotech Incubator as anchors.

## 5. PROJECT MANAGEMENT AND COLLABORATION PLAN

## 6. BROADER IMPACTS

**FIXME:** Describe broader impacts: environmental benefits of energy savings and quality of life benefits to building occupants.

At the undergraduate education level, this work is closely related to **FIXME: describe CSE 132 connection.**

At the graduate education level, this work will support 4 graduate students at Washington Univ. in St. Louis. **FIXME: Expand, including REUs, multidisciplinary angle.**

We will leverage a pair of existing university programs to help us attract students from traditionally underrepresented groups. The Olin Fellowship Program (for women) and the Chancellor's Fellowship Program (aimed at underrepresented minority students) have had a successful track record of enabling individuals to pursue graduate study. In our experience, the most effective method for attracting students from underrepresented groups is by personal contact with a suitable role model. To facilitate this, we regularly ask the appropriately qualified individuals in our group to be actively involved in the recruiting process. This cohort currently includes two minority graduate students (one African-American student and one hispanic student). **FIXME: Can we strengthen the BPC story? Maybe somehow with 132 and maker spaces?**

## 7. RESULTS FROM PRIOR NSF SUPPORT

**CSR: Small: Concurrent Accelerated Data Integration (CNS-1527510, PI: R. Chamberlain), 10/2015–9/2019, \$519,275.**

**Intellectual Merit** – This project investigates the accelerated execution of data integration workflows, which increasingly are bottlenecks in data science. Execution platforms being targeted include both graphics engines and FPGAs. Publications resulting from this work include [8, 9, 22, 28].

**Broader Impacts** – This research project has supported 3 graduate students and 4 REU students. The applications investigated come from the fields of computational biology, astrophysics, and the Internet of Things, further expanding the scope of the students' experience. A benchmark suite of these workflows has been released as a community resource [7].

**TBD - CyberMechProject (CNS-, PI: C. Gill), dates and dollars TBD.**

**Intellectual Merit** – TBD

**Broader Impacts** – TBD

## REFERENCES

- [1] N. Achour and K. Braikia. An MDP-based approach oriented optimal policy for path planning. In *Proc. of International Conference on Machine and Web Intelligence*, pages 205–210, Oct. 2010.
- [2] O. Alagoz, H. Hsu, A. J. Schaefer, and M. S. Roberts. Markov Decision Processes: A Tool for Sequential Decision Making Under Uncertainty. *Medical Decision Making*, 30(4):474–483, 2010.
- [3] M. Alrubaih, M. Zain, M. Alghoul, N. Ibrahim, M. Shameri, and O. Elayeb. Research and development on aspects of daylighting fundamentals. *Renewable and Sustainable Energy Reviews*, 21:494–505, Feb. 2013.
- [4] A. Briggs and M. Sculpher. An Introduction to Markov Modelling for Economic Evaluation. *Pharmacoeconomics*, 13(4):397–409, 1998.
- [5] B. Bueno, J. Wienold, A. Katsifarakis, and T. E. Kuhn. Fener: A radiance-based modelling approach to assess the thermal and daylighting performance of complex fenestration systems in office spaces. *Energy and Buildings*, 94:10–20, Feb. 2015.
- [6] C. Buratti, E. Moretti, E. Bellone, and F. Cotana. Unsteady simulation of energy performance and thermal comfort in non-residential buildings. *Building and Environment*, 59:482–491, Jan. 2013.
- [7] A. M. Cabrera, C. J. Faber, K. Cepeda, R. Derber, C. Epstein, J. Zheng, R. K. Cytron, and R. D. Chamberlain. Data Integration Benchmark Suite v1. DOI: <http://dx.doi.org/10.7936/K7NZ8715>, Apr. 2018.
- [8] A. M. Cabrera, C. J. Faber, K. Cepeda, R. Derber, C. Epstein, J. Zheng, R. K. Cytron, and R. D. Chamberlain. DIBS: A data integration benchmark suite. In *Proc. of ACM/SPIE Int’l Conf. on Performance Engineering Companion*, pages 25–28, Apr. 2018.
- [9] R. D. Chamberlain. Assessing user preferences in programming language design. In *Proc. ACM Int’l Symp. on New Ideas, New Paradigms, and Reflections on Programming and Software*, pages 18–29, Oct. 2017.
- [10] R. D. Chamberlain, C. Ahrens, and C. Gill. Abstractions for cyber-physical systems development: An international opportunity. In *Visioning Workshop for International Networks to Advance CPS Research, Development, and Education Worldwide*, Apr. 2018. Available at <https://cps-vo.org/node/48624>.
- [11] A. Galatioto and M. Beccali. Aspects and issues of daylighting assessment: A review study. *Renewable and Sustainable Energy Reviews*, 66:852–860, Sept. 2016.
- [12] R. Glaubius, T. Tidwell, C. Gill, and W. Smart. ”scheduling policy design for autonomic systems”. *International Journal of Autonomous and Adaptive Communication Systems*, 2(3):276–296, June 2009.
- [13] R. Glaubius, T. Tidwell, C. Gill, and W. Smart. ”real-time scheduling via reinforcement learning”. In *26th Conference on Uncertainty in Artificial Intelligence (UAI)*, July 2010.
- [14] R. Glaubius, T. Tidwell, B. Sidoti, D. Pilla, J. Meden, C. Gill, and W. Smart. ”scalable scheduling policy design for open soft real-time systems”. In *IEEE Real-Time and Embedded Technology and Applications Symposium (RTAS)*, Apr. 2010.
- [15] R. Glaubius, T. Tidwell, W. Smart, and C. Gill. Scheduling design and verification for open soft real-time systems. In *Proc. of Real-Time Systems Symposium*, pages 505–514, Nov. 2008.
- [16] M. Kolonko and H. Benzing. The sequential design of Bernoulli experiments including switching costs. *Operations Research*, 33(2):412–426, Apr. 1985.



- [17] M. Konstantoglou and A. Tsangrassoulis. Dynamic operation of daylighting and shading systems: A literature review. *Renewable and Sustainable Energy Reviews*, 60:268–283, Feb. 2016.
- [18] M. Krarti, P. M. Erickson, and T. C. Hillman. A simplified method to estimate energy savings of artificial lighting use from daylighting. *Building and Environment*, 40:747–754, 2005.
- [19] A. R. Kristensen. A general software system for Markov decision processes in herd management applications. *Computers and Electronics in Agriculture*, 38(3):199–215, Mar. 2003.
- [20] R. Leslie. Capturing the daylight dividend in buildings: why and how? *Building and Environment*, 38:381–385, 2003.
- [21] D. H. W. Li and J. C. Lam. Evaluation of lighting performance in office buildings with daylighting controls. *Energy and Buildings*, 33:793–803, 2001.
- [22] J. Meier, C. Gill, and R. D. Chamberlain. Combining admission and modulation decisions for wireless embedded systems. In *Proc. IEEE 19th Int’l Symp. Real-Time Distributed Computing*, pages 69–78, May 2016.
- [23] J. Meier, S. Sistla, B. Karaus, C. Gill, R. Chamberlain, and T. Tidwell. 16th acm international conference on modeling, analysis and simulation of wireless and mobile systems (mswim). In *”Assessing the Appropriateness of using Markov Decision Processes for RF Spectrum Management”*, Nov. 2013.
- [24] L. Pérez-Lombard, J. Ortiz, and C. Pout. A review on buildings energy consumption information. *Energy and Buildings*, 40:394–398, 2008.
- [25] P. Plympton, S. Conway, and K. Epstein. Daylighting in schools: Improving student performance and health at a price schools can afford. In *Proc. of American Solar Energy Society Conference*, June 2000.
- [26] M. L. Puterman. *Markov Decision Processes*. John Wiley & Sons, Inc., Hoboken, NJ, 1994.
- [27] B. Schaban-Maurer. *Rise of the Citizen Practitioner: A Phronesis-Based Approach to Citizen Engagement and Social Policy*. Scholar’s Press, 2013.
- [28] J. A. Shidal. *Exploiting the Weak Generational Hypothesis for Write Reduction and Object Recycling*. PhD thesis, Dept. of Computer Science and Engineering, Washington University in St. Louis, May 2016.
- [29] T. Tidwell, C. Bass, E. Lasker, M. Wylde, C. Gill, and W. Smart. ”scalable utility aware scheduling heuristics for real-time tasks with stochastic non-preemptive execution intervals”. In *23rd Euromicro Conference on Real-Time Systems (ECRTS)*, pages 238–247, July 2011.
- [30] T. Tidwell, R. Glaubius, C. D. Gill, and W. D. Smart. Optimizing expected time utility in cyber-physical systems schedulers. In *Proc. of IEEE 31st Real-Time Systems Symposium*, pages 193–201, Dec. 2010.
- [31] D. Vázquez-Moliní, M. González-Montes, A. Álvarez Fernández-Balbuena, Ángel García-Botella, W. Pohl, T. Galan, and E. Bernabéu. Horizontal daylighting system for office buildings. *Energy and Buildings*, 67:525–530, 2013.
- [32] D. J. White. A Survey of Applications of Markov Decision Processes. *The Journal of the Operational Research Society*, 44(11):1073–1096, Nov. 1993.
- [33] K. K. Yin, G. G. Yin, and H. Liu. Stochastic modeling for inventory and production planning in the paper industry. *AIChE Journal*, 50(11):2877–2890, Nov. 2004.